

Machine Learning Starter Kit

1) What is Machine Learning?

Machine Learning (ML) is a collection of methods that learn patterns from data to make predictions or decisions without being explicitly programmed for each rule.

- **Input → Model → Output:** Given examples (data), an algorithm fits a *model* that maps inputs to outputs.
- **Goal:** Generalize—perform well on *new, unseen* data, not just the training data.

Everyday examples: photo tag suggestions, spam filtering, product recommendations, speech recognition, fraud detection.

2) Major Types of ML

1. **Supervised Learning** — Learn from labeled examples.
 2. **Tasks:** **Classification** (discrete labels), **Regression** (continuous values)
 3. **Examples:** Digit recognition, house price prediction
 4. **Unsupervised Learning** — Find structure in unlabeled data.
 5. **Tasks:** **Clustering** (k-means), **Dimensionality reduction** (PCA)
 6. **Examples:** Customer segmentation, anomaly detection
 7. **Semi-Supervised Learning** — Mix of labeled + unlabeled data to improve performance when labels are scarce.
 8. **Self-Supervised Learning** — Create labels from the data itself (e.g., predicting masked words/images); often used to pretrain large models.
 9. **Reinforcement Learning (RL)** — Learn by trial and error with rewards (games, robotics, operations).
 10. **Other distinctions:**
 11. **Online vs. Batch** learning (streaming updates vs. periodic retraining)
 12. **Generative vs. Discriminative** models
-

3) End-to-End ML Workflow (at a glance)

```
Data → Split → Preprocess/Normalize → Train → Validate/Evaluate → Serialize →  
Load & Predict → Monitor
```

- **Data acquisition:** collect/clean data; record provenance and licenses.
- **Split:** hold-out test set (and optionally validation set) *before* peeking at results.
- **Preprocess/Normalize:** handle missing values, scale features, encode categories.
- **Train:** fit the model on the training set.

- **Validate/Evaluate:** tune hyperparameters; measure generalization.
- **Serialize:** save the trained model artifact (e.g., `joblib` file).
- **Load & Predict:** reload model for inference; ensure same preprocessing.
- **Monitor:** check for drift; retrain as needed.

Good practice defaults

- Keep a **strict test set** (never used in training or tuning).
 - Use **cross-validation** for robust estimates when data is limited.
 - Prefer **simple baselines first** (logistic regression, linear models).
 - Track **metrics** and **random seeds** for reproducibility.
 - Save the **preprocessing steps** together with the model or document them clearly.
-

4) Essential Concepts & Glossary

- **Feature:** Measurable property used by the model.
 - **Label/Target:** The thing you want to predict.
 - **Sample/Instance:** One row/example in your dataset.
 - **Overfitting:** Model memorizes training data, performs poorly on new data.
 - **Underfitting:** Model is too simple; misses important patterns.
 - **Bias/Variance:** Systematic error vs. sensitivity to data fluctuations.
 - **Regularization:** Techniques (e.g., L2) to discourage overly complex models.
 - **Hyperparameters:** Settings chosen before training (e.g., learning rate, C).
 - **Cross-Validation (CV):** Repeated train/validation splits to estimate performance.
 - **Normalization/Standardization:** Scale features (e.g., zero mean, unit variance).
 - **Confusion Matrix:** Table of TP/FP/TN/FN for classification.
 - **Precision/Recall/F1:** Metrics focusing on positive predictions and coverage.
 - **ROC-AUC/PR-AUC:** Threshold-independent metrics for classifiers.
 - **MSE/RMSE/MAE:** Regression error metrics.
 - **Pipeline:** Chaining preprocessing and model steps together.
 - **Serialization:** Saving model/preprocessing to disk to reuse later.
-

5) Popular Datasets & Sources

- **Kaggle** (competitions, curated datasets, notebooks). Useful for realistic projects. Requires free account and API token for programmatic download.
- **Hugging Face Hub** (datasets, models). Strong coverage for text, vision, audio; easy programmatic access.
- **UCI Machine Learning Repository** (classic academic datasets).
- **scikit-learn built-ins** (e.g., digits, iris, wine): great for teaching.

Tip: Start on built-in datasets (no auth, tiny downloads), then move to Kaggle/Hugging Face to practice real-world data handling.

6) Evaluation & Experimentation

- **Train/Validation/Test:** 60/20/20 or 70/15/15 are common starting points.
 - **Baselines:** Always compare to a simple baseline (e.g., majority class) and a simple model (e.g., logistic regression).
 - **Model selection:** Use validation or CV—not the test set—to choose models.
 - **Reporting:** Include metric definitions, chosen splits, hyperparameters, and random seeds.
-

7) Common Pitfalls & How to Avoid Them

- **Data leakage:** Accidentally using test information during training. *Fix:* split early; fit scalers/encoders only on training data.
 - **Mismatch preprocessing:** Different preprocessing at train vs. predict time. *Fix:* save preprocessors or use a single pipeline.
 - **Imbalanced classes:** Accuracy can be misleading. *Fix:* use precision/recall/F1, stratified splits, resampling, or class weights.
 - **Non-reproducible results:** Randomness and ad hoc code. *Fix:* set seeds, fix versions, log experiments.
-

8) Minimal Hands-On Example (Digits Classification)

Goal: Recognize hand-written digits using a simple, readable pipeline in Python.

Steps covered: acquire → split → normalize → train → validate → serialize → load/predict.

Tools: `scikit-learn` (datasets, model, metrics), `joblib` (save/load). Minimal imports, no heavy frameworks.

See the provided `digit_pipeline.py` file for a fully runnable script with comments your students can follow line-by-line.

9) Where Kaggle & Hugging Face Fit In (for Data Collection)

- **Kaggle:**

- Create an account → Account settings → Create API token.
- Install `kaggle` CLI locally, place `kaggle.json` in `~/.kaggle/`.
- `kaggle datasets download -d <owner/dataset>` to fetch zip files.
- Unzip, inspect README/licenses, then proceed with your pipeline.

- **Hugging Face Datasets:**

- `pip install datasets`

- `from datasets import load_dataset` → `load_dataset("mnist")`, etc.

- Convert to NumPy/Pandas, then follow the same pipeline steps.

The project starter file (`ml_projects_starter.py`) includes optional snippets and TODOs for both sources.

10) Training Process: Parameters, Epochs & Optimization

10.1 Core terminology

- **Loss function:** A number the optimizer tries to *minimize* (e.g., cross-entropy for classification, MSE for regression).
- **Parameters:** Values learned from data (e.g., weights/coefficients).
- **Hyperparameters:** Settings you choose before training (e.g., learning rate, `C`, depth, batch size, number of epochs).
- **Epoch:** One full pass over the training set.
- **Batch (minibatch):** A small subset of the data used to compute an update.
- **Iteration / Step:** One optimizer update (one batch processed).
- **Optimization:** The procedure that adjusts parameters to reduce loss (e.g., gradient descent variants).
- **Convergence:** When further updates change the loss/parameters only slightly (meeting a tolerance).

In **scikit-learn**: many solvers (e.g., `lbfgs` in `LogisticRegression`) use *iterations* with a stopping **tolerance** (`tol`) and **max_iter**. `SGDClassifier/Regressor` expose **epochs**, **batch size**, and **learning rate** more directly.

10.2 Common optimization algorithms

- **Batch Gradient Descent:** Uses the whole dataset per step (stable but slow).
- **Stochastic Gradient Descent (SGD):** Uses one sample per step (fast but noisy).
- **Mini-batch SGD:** Uses small batches (common default; good trade-off).
- **Momentum / Nesterov:** Adds velocity to smooth updates and accelerate.
- **Adaptive methods (Adam/Adagrad/RMSProp):** Scale learning rates per parameter (fast, often good defaults for deep nets).

10.3 Key hyperparameters (what they mean & typical effects)

- **Learning rate (η):** Size of each update.
- Too high → divergence/instability; too low → slow/underfitting.
- Often the *most important* hyperparameter to tune.

- **Batch size:**
 - Small batches → noisier gradients, better generalization sometimes.
 - Large batches → faster per-epoch on GPUs, may need **learning-rate warm-up**.
 - **Epochs / max_iter:**
 - Too few → underfitting; too many → overfitting unless regularized/early stopped.
 - **Regularization strength** (e.g., C in logistic regression, α / λ elsewhere):
 - Higher regularization → simpler models, less overfitting, possibly lower training accuracy but better test performance.
 - **Tolerance (ϵ) & early-stopping criteria:**
 - Stop when improvement falls below a threshold for N checks.
 - **Model-specific:**
 - Trees/forests: depth, number of estimators, min samples per split/leaf.
 - SVM: kernel, C , γ .
 - kNN: number of neighbors k , distance metric.
-

10.4 Preventing overfitting (regularization & friends)

- **L2 / L1 regularization:** Penalize large weights; L1 can induce sparsity.
- **Early stopping:** Monitor validation loss; stop when it stops improving.
- **Data augmentation** (vision/text/audio): Create realistic variations of inputs.
- **Dropout** (deep learning): Randomly zero some activations during training.
- **Ensembling:** Combine models (bagging, boosting, stacking).
- **Proper split:** Keep a *strict* validation/test set; avoid data leakage.

For **linear models & SVMs, feature scaling** (standardization) is crucial. For **tree-based models**, scaling is usually unnecessary.

10.5 Learning-rate schedules (when training for many epochs)

- **Step decay:** Reduce η by a factor at fixed epochs.
 - **Exponential/cosine decay:** Smoothly reduce η over time.
 - **Warm-up:** Start with smaller η for a few epochs, then increase to target η .
 - **Cyclical / One-cycle:** Vary η within ranges to escape sharp minima.
-

10.6 Choosing and monitoring metrics

- **Classification:** Accuracy, precision/recall/F1, ROC-AUC, PR-AUC; inspect the **confusion matrix**.
 - **Regression:** MAE (robust), RMSE (penalizes large errors), R^2 .
 - Track **training vs. validation** curves to spot under/overfitting:
 - Both low → underfitting (increase model capacity, train longer, reduce reg).
 - Train low / Val high → overfitting (more reg, more data/augmentation, early stopping).
-

10.7 A practical tuning workflow (checklist)

1. **Start simple:** Small model + reasonable defaults; verify data pipeline & splits.
 2. **Scale features** (if needed) and set a **reproducible seed**.
 3. **Pick a baseline metric** and log it.
 4. **Tune learning rate** (and batch size if applicable) first.
 5. **Adjust model capacity / regularization** next.
 6. **Use validation curves** or **cross-validation** for robust estimates.
 7. **Introduce schedules/early stopping** to stabilize training.
 8. **Lock the model** and evaluate once on the **held-out test set**.
 9. **Serialize** the exact preprocessing + model; record versions, seeds, and hyperparameters.
-

10.8 Hyperparameter search strategies

- **Manual search:** Fast to start; good when you know what matters.
- **Grid search:** Systematic but can be wasteful in high-dimensional spaces.
- **Random search:** Surprisingly strong; explores more configurations under the same budget.
- **Bayesian / adaptive search** (e.g., TPE, Gaussian processes): Model the performance surface to choose promising trials.

In **scikit-learn**: use `GridSearchCV` / `RandomizedSearchCV` with **pipelines** so preprocessing is fit only on training folds.

10.9 Troubleshooting guide

- **Training loss not decreasing:** Lower learning rate; check data/labels; verify scaling; increase `max_iter`.
 - **Validation worse than training:** Add regularization, early stopping, augmentation; collect more data.
 - **Model is slow:** Reduce batch size memory peaks (DL); simplify model; subsample features; use fewer estimators.
 - **Unstable results:** Set random seeds; fix library versions; increase dataset size or use cross-validation.
-

10.10 Mapping to the digits example (logistic regression)

- Relevant knobs: `max_iter`, `tol`, **regularization** (`C`, `penalty`), and **scaler** choice.
- Typical path:
- Standardize features (already in pipeline).
- Increase `max_iter` until convergence; adjust `tol` if needed.
- Tune `C` on the validation set or via `GridSearchCV`.
- Serialize both **scaler + model**; reload to predict consistently.

11) Further Reading

- *Hands-On Machine Learning with Scikit-Learn, Keras & TensorFlow* (Aurélien Géron)
- *The Hundred-Page Machine Learning Book* (Andriy Burkov)
- scikit-learn User Guide & Tutorials (official)
- Hugging Face Datasets documentation
- Kaggle Learn micro-courses