

# Machine Learning Algorithms & Deep Learning - Lecture + Cheatsheet

**How training actually works.** In supervised learning we observe inputs and labeled targets and learn a function that generalizes. You choose a **model family**, define a **loss**, and use an **optimizer** to update **parameters**. We monitor a validation set to avoid **overfitting** and finally **serialize** the trained model for reuse.

- **Parameters:** values learned from data (weights, biases).
- **Hyperparameters:** settings you decide (learning rate, batch size, epochs, regularization, depth).
- **Loss:** error measure (MSE, cross-entropy).
- **Optimizer:** update rule (Gradient Descent, SGD, Adam).

**Epochs, batches, iterations.** One **epoch** scans the whole training set. We split data into **mini-batches**; each batch produces a gradient step - an **iteration**. The **batch size** controls the noise/efficiency trade-off. Paired with **learning rate**, it shapes convergence speed and stability.

- **Small batches** -> noisier updates, regularization effect, lower memory.
- **Large batches** -> smoother updates, better hardware utilization.
- **LR too high** -> divergence; **too low** -> slow/plateau; consider warm-up/decay schedules.

## Supervised algorithms - when to pick what

Use this as a quick selection guide; each item includes one-line intuition + key knobs.

- **Linear regression:** continuous targets; add **L2/L1** for stability/feature selection.
- **Logistic regression:** strong linear baseline for classification; interpretable.
- **k-NN:** prediction by neighbors; sensitive to scaling; choose **k** via validation.
- **SVM (RBF):** maximum-margin with **non-linear** kernel; tune **C, gamma**.
- **Decision trees:** human-readable rules; prune to curb overfitting.
- **Random Forests:** ensemble of trees; robust default for tabular data.
- **Gradient Boosting** (XGBoost/LightGBM): sequential trees that fix residuals; tune **lr, depth, estimators**.
- **Naive Bayes:** simple probabilistic model; great for bag-of-words text.

## Unsupervised algorithms - structure without labels

- **k-Means:** partitions into k clusters; pick k via elbow/silhouette.
- **Hierarchical clustering:** dendrogram for exploratory analysis.
- **DBSCAN:** density-based clusters and outliers; tune **eps, min\_samples**.
- **PCA:** linear **dimensionality reduction** for visualization/denoising.

## Deep Learning - what it is and how it works

**Definition.** Deep learning is a subfield of ML that learns hierarchical, non-linear representations using multi-layer **neural networks**. Unlike many classical models that rely on hand-engineered features, deep networks learn useful features from raw data end-to-end.

**Neural network workflow - from input to output.**

**1) Neuron (node) computation.** Each neuron takes a vector of inputs and computes a weighted sum plus a bias:  $z = w_1*x_1 + w_2*x_2 + \dots + b$ . It then applies an activation function to produce its output  $a = \text{activation}(z)$ . Common activations are **ReLU** ( $\max(0, z)$ ), **sigmoid** ( $1/(1+e^{-z})$ ), and **tanh**. Activations introduce non-linearity so stacked layers can model complex functions.

**2) What each layer means.**

- **Input layer:** holds raw features (pixels, words, tabular columns). No computation here; it just feeds the next layer.
- **Hidden layers:** each hidden layer transforms its inputs into progressively more abstract features. Earlier layers capture simple patterns; later layers combine them into higher-level concepts.
- **Output layer:** produces final predictions. For classification we often use a softmax layer; for regression a linear output.


**3) Forward pass.** Data flows layer by layer: inputs  $\rightarrow$  hidden layers  $\rightarrow$  outputs. At each layer we compute weighted sums and activations to get the next representation. The result is a prediction  $y_{\hat{}}$ .

**4) Loss.** We compare predictions  $y_{\hat{}}$  with true targets  $y$  using a loss function (cross-entropy for classification, mean squared error for regression). The loss is a single number that summarizes how wrong the network is on the current batch.

**5) Backpropagation.** Using the chain rule, we propagate gradients of the loss backward through the network to compute how much each weight and bias contributed to the error.

**6) Parameter update.** An optimizer such as SGD or Adam updates weights using the gradients. Key dials are **learning rate**, **batch size**, and number of **epochs**.

**7) Repeat over epochs.** We repeat forward pass  $\rightarrow$  loss  $\rightarrow$  backprop  $\rightarrow$  update over many mini-batches and epochs until the validation metric stops improving (early stopping) or we hit a budget.



### Common architectures.

- **MLP (Multilayer Perceptron)**: also called a feed-forward network; good for tabular or generic signals.
- **CNN (Convolutional Neural Network)**: uses convolution filters for images and spatial data; captures local patterns with weight sharing.
- **RNN (Recurrent Neural Network)**: processes sequences by carrying state over time; classic variants are **LSTM (Long Short-Term Memory)** and **GRU (Gated Recurrent Unit)**.
- **Transformer**: based on self-attention; dominant in **NLP (Natural Language Processing)** and increasingly used in vision and audio.

### Training tips.

- **Initialization**: good weight initialization speeds up training and avoids early saturation.
- **Normalization**: batch norm or layer norm helps stabilize deep stacks.
- **Regularization**: dropout, weight decay, data augmentation, early stopping.

**Put it together.** Start with a simple baseline, measure a clear metric, iterate with disciplined validation. Once satisfied, save the model so it is reusable in code and deployment.

- **Workflow**: split (train/val/test) -> train -> validate -> tune -> test once.
- **Search**: coarse grid/random, then refine around good regions.
- **Serialize**: scikit-learn **joblib/pickle**; PyTorch **state\_dict** (.pt/.pth); TensorFlow/Keras **SavedModel**.

## Sources

- Infinite Codes - “All Machine Learning algorithms explained in 17 min” (YouTube, id: E0Hmnixke2g).
- Infinite Codes - “Epochs, Iterations and Batch Size | Deep Learning Basics” (YouTube, id: SftOqbMrGfE).
- scikit-learn Documentation - Model persistence; Hyper-parameter tuning; Cross-validation.